

Processes

Course Code: CSC 2209

Course Title: Operating Systems



Dept. of Computer Science
Faculty of Science and Technology

Lecturer No:	04	Week No:	04	Semester:	
Lecturer:	<i>Name & email</i>				

Lecture Outline



1. Process Concept
2. Process Scheduling
3. Operations on Processes
4. Inter-process Communication (IPC)
5. Examples of IPC Systems
6. Communication in Client-Server Systems


Process Concept

- ❑ **Process** – a program in execution; process execution **must** progress in sequential fashion
- ❑ Program is *passive* entity stored on disk (**executable file**), process is *active*
 - ❑ Program becomes process when executable file loaded into memory
- ❑ Execution of program started via:
 - ❑ GUI mouse clicks,
 - ❑ command line entry of its name,
 - ❑ etc.
- ❑ **One program can be several processes**
 - ❑ Consider multiple users executing the same program

Process Structure


- ❑ A process is more than the program code, which is sometimes known as the **text section**.
- ❑ It also includes the current activity:
 - ❑ The value of the **program counter**
 - ❑ The contents of the **processor's registers**.
- ❑ It also includes the process **stack**, which contains temporary data (such as function parameters, return addresses, and local variables)
- ❑ It also includes the **data section**, which contains global variables, static variable (once created stay forever).
- ❑ It may also include a **heap**, which is memory that is dynamically allocated during process run time.

Process in Memory




Process State

- ❑ As a process executes, it changes **state**
 - ❑ **new**: The process is being created
 - ❑ **ready**: The process is waiting to be assigned to a processor
 - ❑ **running**: Instructions are being executed
 - ❑ **waiting**: The process is waiting for some event to occur
 - ❑ **terminated**: The process has finished execution



Process State

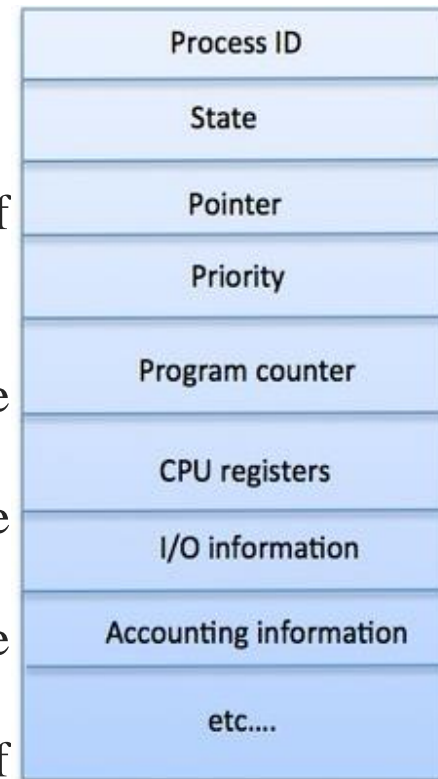
□ Diagram of Process State



Process Control Block (PCB)

Information associated with each process (use to present a process in OS) (also called **task control block**)

- ☐ Process ID / Number – Unique ID.
- ☐ Process state – running, waiting, etc.
- ☐ Program counter – location of instruction to next execute (addr of next inst)
- ☐ CPU registers – contents of all process-centric registers ()
- ☐ CPU scheduling information- priorities, scheduling queue pointers
- ☐ Memory-management information – memory allocated to the process
- ☐ Accounting information – CPU used, clock time elapsed since start, time limits
- ☐ I/O status information – I/O devices allocated to process, list of open files




Threads

- ❑ So far, process has a single thread of execution
- ❑ Consider having multiple program counters per process
 - ❑ Multiple locations can execute at once
 - ❑ Multiple threads of control -> **threads**
- ❑ Need storage for thread details, multiple program counters in PCB
- ❑ Covered in the next chapter

Process Representation in Linux

Represented by the C structure `task_struct`

```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice; /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```

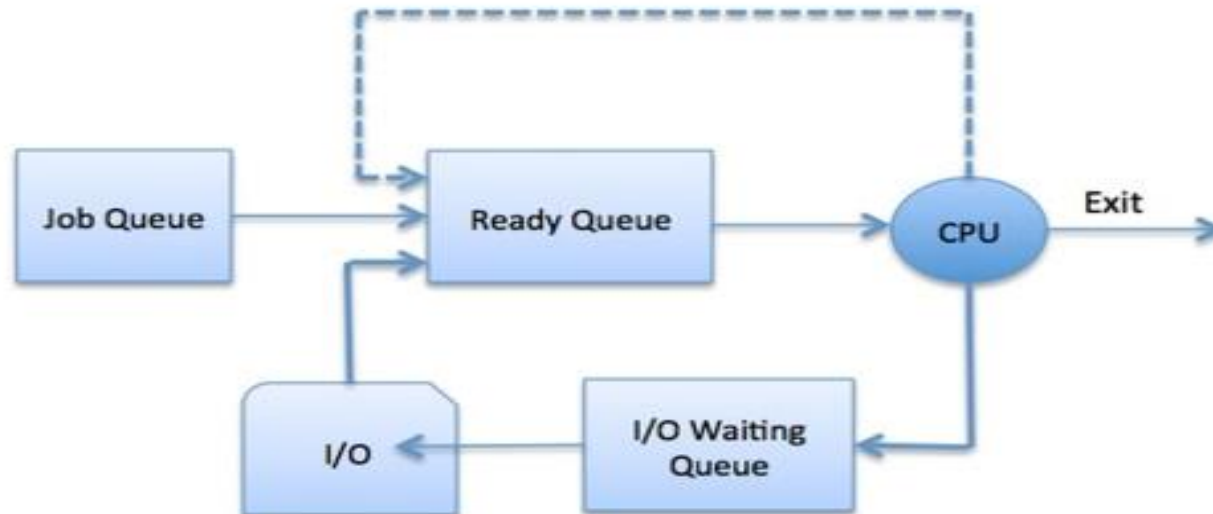


Process Scheduling

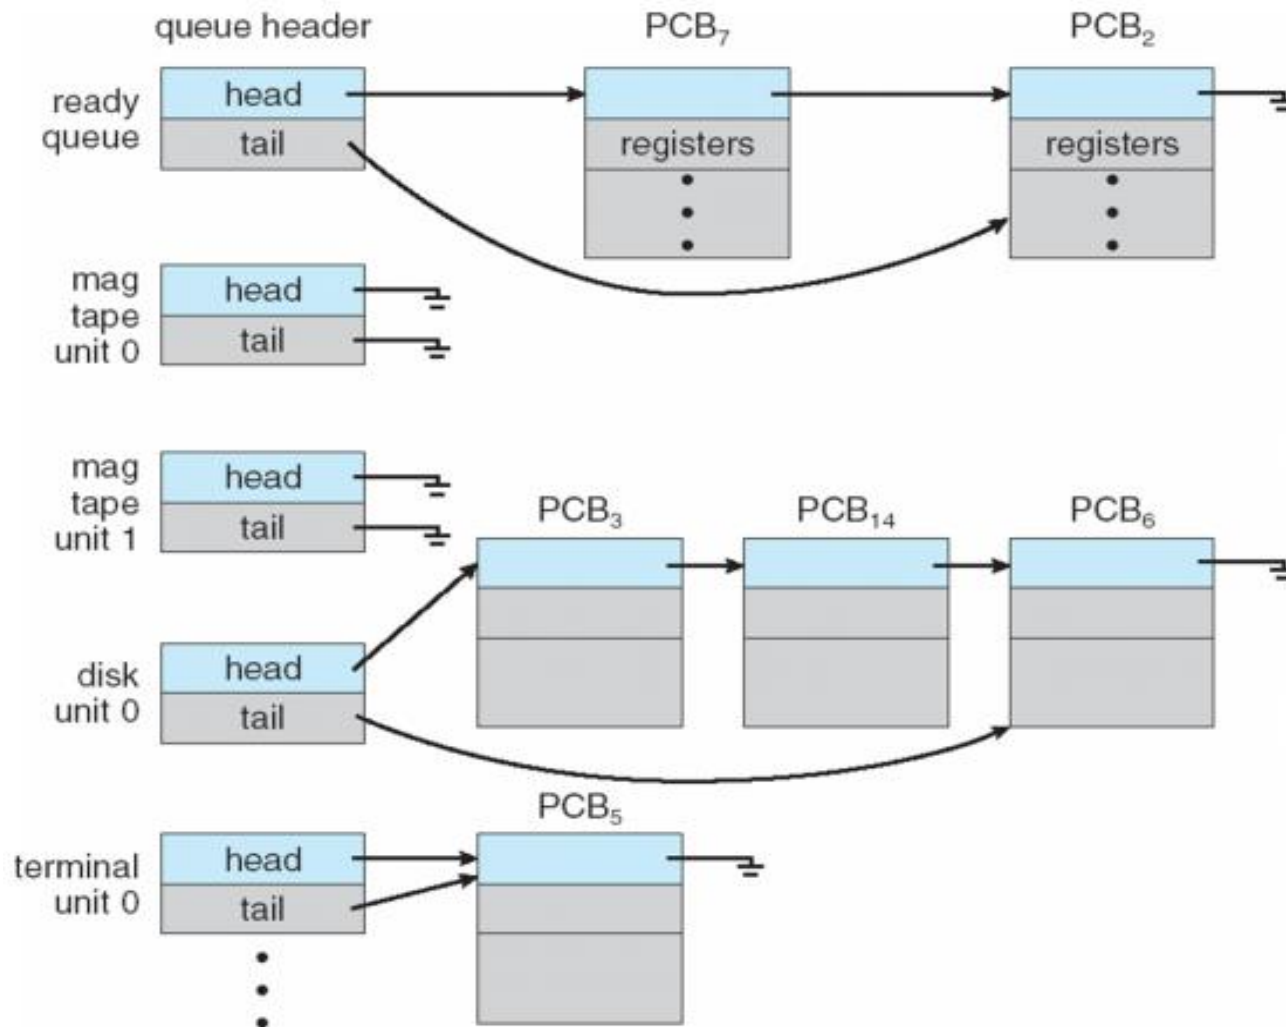
- ❑ Maximize CPU use
 - ❑ Quickly switch processes onto CPU for time sharing
- ❑ Process “gives” up the CPU under two conditions:
 - ❑ I/O request
 - ❑ After N units of time have elapsed (need a timer)
- ❑ Once a process gives up the CPU it is added to the “ready queue”
- ❑ **Process scheduler** selects among available processes in the ready queue for next execution on CPU

Scheduling Queues

- ❑ OS Maintains **scheduling queues** of processes
 - ❑ **Job queue** – set of all processes in the system
 - ❑ **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - ❑ **Device queues** – set of processes waiting for an I/O device
 - ❑ Processes migrate among the various queues

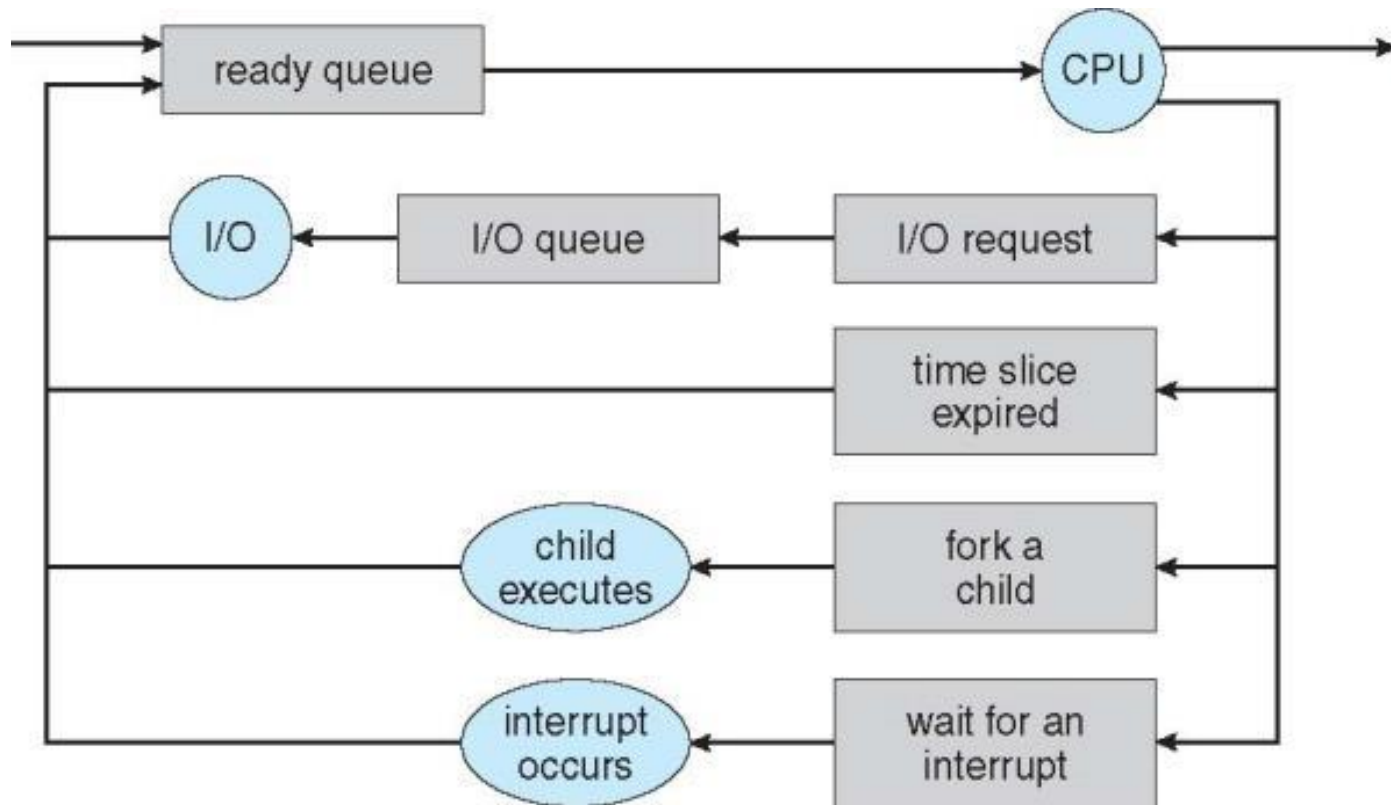


Ready Queue And Various I/O Device Queues

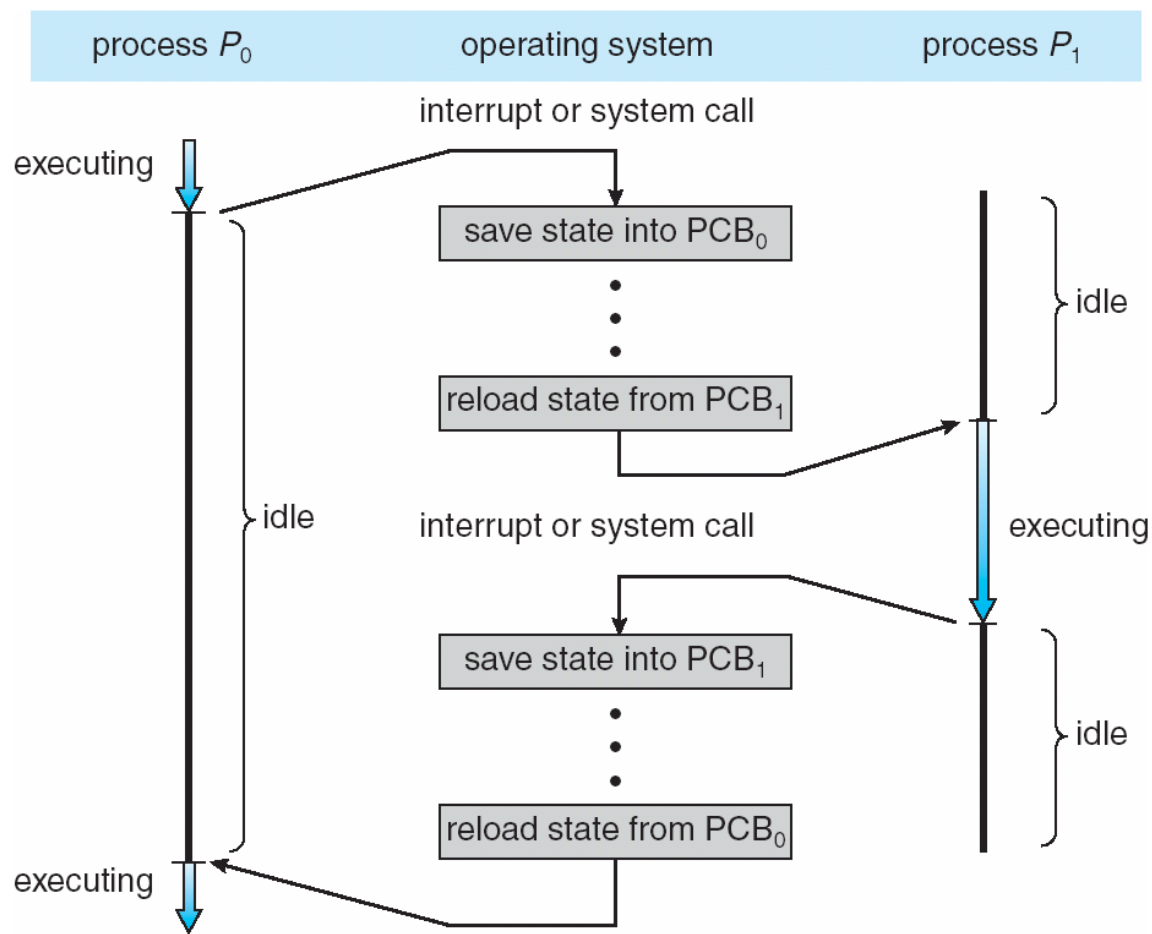


Representation of Process Scheduling

- ❑ **Queuing diagram** represents queues, resources, flows



CPU Switch From Process to Process



Schedulers

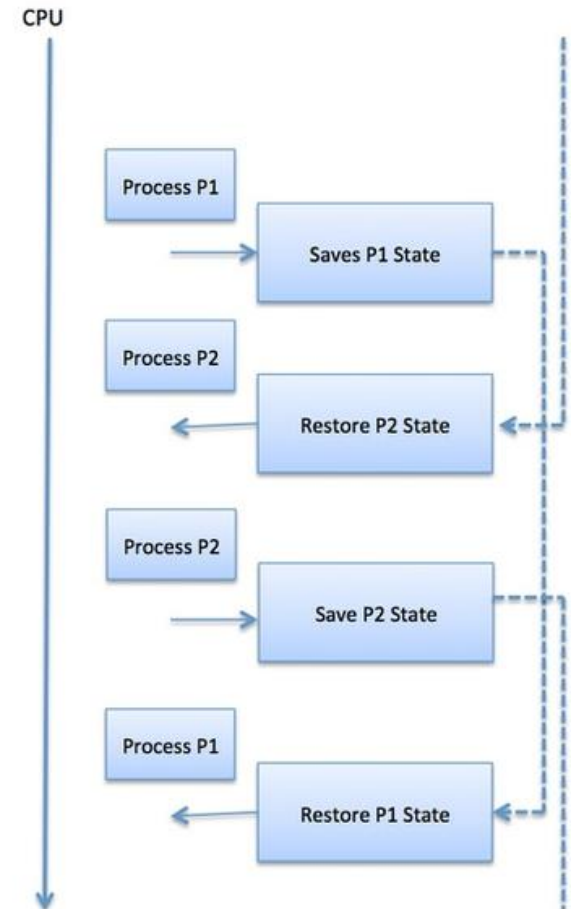
- ❑ **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates a CPU
 - ❑ Sometimes the only scheduler in a system
 - ❑ Short-term scheduler is invoked frequently (milliseconds) \Rightarrow (must be fast)
- ❑ **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
 - ❑ Long-term scheduler is invoked infrequently (seconds, minutes) \Rightarrow (may be slow)
 - ❑ The long-term scheduler controls the **degree of multiprogramming**
- ❑ Processes can be described as either:
 - ❑ **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - ❑ **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- ❑ Long-term scheduler strives for good *process mix*

Multitasking in Mobile Systems

- ❑ Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended
- ❑ Starting with **iOS 4**, it provides for a
 - ❑ Single **foreground** process – controlled via user interface
 - ❑ Multiple **background** processes – in memory, running, but not on the display, and with limits
 - ❑ Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- ❑ **Android** runs foreground and background, with fewer limits
 - ❑ Background process uses a **service** to perform tasks
 - ❑ **Service** can keep **running even if background process is suspended**
 - ❑ **Service has no user interface, small memory use**

Context Switch

- ❑ When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- ❑ **Context** of a process represented in the PCB
- ❑ **Context-switch time is pure overhead**; the system does no useful work while switching
 - ❑ The more complex the OS and the PCB → the longer the context switch
- ❑ Time dependent on hardware support
 - ❑ Some hardware provides multiple sets of registers per CPU → multiple contexts loaded at once




Operations on Processes

- ☐ System must provide mechanisms for:
 - ☐ process creation,
 - ☐ process termination,
 - ☐ and so on as detailed next

Process Creation

- ❑ A **process** may create other processes.
- ❑ **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- ❑ Generally, a process is identified and managed via a **process identifier (pid)**
- ❑ A Tree of Processes in UNIX




Process Creation (cont'd)

- ☐ Resource sharing among parents and children options
 - ☐ Parent and children share all resources
 - ☐ Children share subset of parent's resources
 - ☐ Parent and child share no resources
- ☐ Execution options
 - ☐ Parent and children execute concurrently
 - ☐ Parent waits until children terminate

Process Creation (cont'd)

- ❑ Address space
 - ❑ A child is a duplicate of the parent address space.
 - ❑ A child loads a program into the address space.
- ❑ UNIX examples
 - ❑ **fork()** system call creates new process
 - ❑ **exec()** system call used after a **fork()** replaces the process' memory space with a new program



C program to create a separate process in UNIX

```
int main()
{
    pid_t pid;
    /*fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /*child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}
```

C Program Forking Separate Process

- ❑ The C program illustrates how to create a new process UNIX.
- ❑ After **fork()** there are **two different processes running copies of the same program**. The only difference is that the value of `pid` for the child process is zero, while that for the parent is an integer value greater than zero
- ❑ The child process inherits privileges and scheduling attributes from the parent, as well certain resources, such as open files.
- ❑ The child process then overlays its address space with the UNIX command “ls” (used to get a directory listing) using the `execlp()` (a version of the `exec()` system call).
- ❑ The parent waits for the child process to complete with the `wait()` system call.
- ❑ When the child process completes, the parent process resumes from the call to `wait()`, where it completes using the `exit()` system call.

Creating a Separate Process via Windows API

```
int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
```

Creating a Separate Process via Windows API (cont'd)

```
{
fprintf(stderr, "Create Process Failed");
return -1;
}
/* parent will wait for the child to complete */
WaitForSingleObject(pi.hProcess, INFINITE);
printf("Child Complete");

/* close handles */
CloseHandle(pi.hProcess);
CloseHandle(pi.hThread);
}
```

Process Termination

- ❑ A process terminates when it finishes executing its final statement and it asks the operating system to delete it by using the **exit() system call**.
 - ❑ At that point, the process may return a status value (typically an integer) to its parent process (via the wait() system call).
 - ❑ All the resources of the process are deallocated by the operating system.
- ❑ A parent may terminate the execution of children processes using the **abort() system call**. Some reasons for doing so:
 - ❑ Child has exceeded allocated resources
 - ❑ Task assigned to child is no longer required
 - ❑ The parent is exiting and the operating systems does not allow a child to continue if its parent terminates

Process Termination (cont'd)

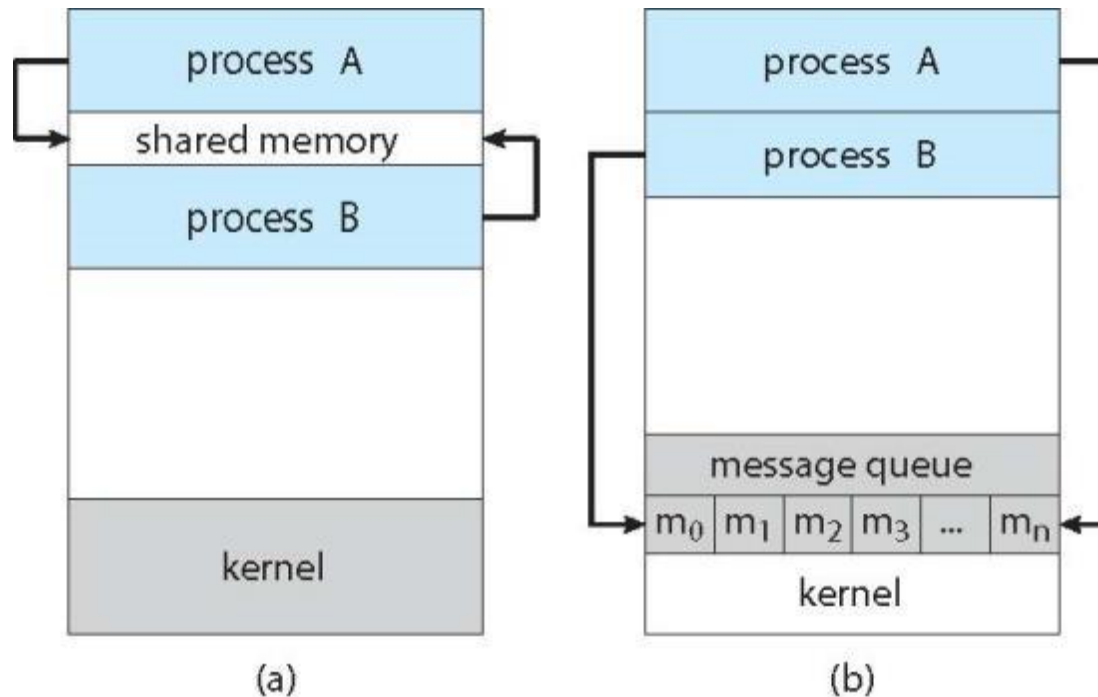
- ❑ Some operating systems do not allow a child process to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.
 - ❑ **cascading termination.** All children, grandchildren, etc. are terminated.
 - ❑ The termination is initiated by the operating system.
- ❑ The **parent process** may wait for termination of a child process by using the **wait()** system call. The call returns status information and the pid of the terminated process
- ❑ **pid = wait(&status);**
- ❑ If no parent waiting (did not invoke **wait()**) process is **zombie**
 - ❑ **Status table yes but running no but parent process need to call exit() system**
- ❑ If parent terminated without invoking **wait**, process is **orphan**
- ❑ **Running yes but parent running no**

Inter process Communication (IPC)

- ❑ Processes within a system may be *independent* or *cooperating*
 - ❑ **Cooperating processes** can affect or be affected by other processes, including sharing data
 - ❑ **Independent processes** cannot affect other processes
- ❑ Reasons for having **cooperating** processes:
 - ❑ Information sharing
 - ❑ Computation speedup (multiple processes running in parallel- subtasks divided into processes)
 - ❑ Modularity (system design- several modules/process need comm.)
 - ❑ Convenience (multitasking-multiprogramming make user life easy)
- ❑ Cooperating processes need **inter process communication (IPC)**

Communications Models

- ❑ Two models of IPC (Inter process communication)
 - ❑ **Shared memory**
 - ❑ **Message passing**



Shared Memory Systems

- ❑ **A region of memory that is shared by cooperating processes.**
- ❑ An area of memory shared among the processes that wish to communicate
- ❑ The communication is under the control of the users processes not the operating system.
- ❑ Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.

Synchronization

- ❑ Cooperating processes that access shared data need to synchronize their actions to ensure data consistency
- ❑ Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- ❑ Illustration of the problem – The producer-Consumer problem
 - ❑ Producer process produces information that is consumed by a Consumer process.
 - ❑ The information is passed from the Producer to the Consumer via a buffer.
 - ❑ Two types of buffers can be used:
 - ❑ **unbounded-buffer** places no practical limit on the size of the buffer
 - ❑ **bounded-buffer** assumes that a fixed buffer size

Bounded-Buffer Solution

- ❑ Shared data

- ❑ `#define BUFFER_SIZE 10`
 - ❑ `typedef struct {`
 - ❑ `. . .`
 - ❑ `} item;`

 - ❑ `item buffer[BUFFER_SIZE];`
 - ❑ `int in = 0;`
 - ❑ `int out = 0;`

- ❑ Solution presented in the next two slides is correct, but only 9 out of 10 buffer elements can be used

Bounded-Buffer – Producer

```
item next_produced;
while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Bounded Buffer – Consumer

```
item next_consumed;

while (true) {
    while (in == out)

        ; /* do nothing */
    next_consumed = buffer[out];

    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next consumed */

}
```

Message Passing Systems

- ❑ Mechanism for processes to communicate and to synchronize their actions (messages are exchanged between processes)
 - ❑ Without resorting to shared variables
- ❑ IPC facility provides two operations:
 - ❑ **send**(*message*)
 - ❑ **receive**(*message*)
- ❑ The *message* size is either **fixed or variable**

Message Passing (cont'd)

- ❑ If processes P and Q wish to communicate, they need to:
 - ❑ Establish a **communication link** between them
 - ❑ Exchange messages via send/receive
- ❑ Implementation issues:
 - ❑ How are links established?
 - ❑ Can a link be associated with more than two processes?
 - ❑ How many links can there be between every pair of communicating processes?
 - ❑ What is the capacity of a link?
 - ❑ Is the size of a message that the link can accommodate fixed or variable?
 - ❑ Is a link unidirectional or bi-directional?

Implementation of Communication Link

➤ Physical:

- Shared memory
- Hardware bus
- Network

➤ Logical:

- Direct or indirect
- Synchronous or asynchronous
- Automatic or explicit buffering

Direct Communication

- ❑ Processes must name each other explicitly:
 - ❑ **send** (*P*, *message*) – send a message to process *P*
 - ❑ **receive**(*Q*, *message*) – receive a message from process *Q*
- ❑ Properties of communication link
 - ❑ Links are established automatically
 - ❑ A link is associated with exactly one pair of communicating processes
 - ❑ Between each pair there exists exactly one link
 - ❑ The link may be unidirectional, but is usually bi-directional

Indirect Communication

- ❑ Messages are directed and received from mailboxes (also referred to as ports)
 - ❑ Each mailbox has a unique id
 - ❑ Processes can communicate only if they share a mailbox
- ❑ Operations
 - ❑ create a new mailbox (port)
 - ❑ send and receive messages through mailbox
 - ❑ delete a mailbox
- ❑ Primitives are defined as:
 - ❑ **send**(*A, message*) – send a message to mailbox A
 - ❑ **receive**(*A, message*) – receive a message from mailbox A

Indirect Communication (cont'd)

- ❑ Properties of communication link
 - ❑ Link established only if processes share a common mailbox
 - ❑ A link may be associated with many processes
 - ❑ Each pair of processes may share several communication links
 - ❑ Link may be unidirectional or bi-directional

Indirect Communication Issues

- ❑ Mailbox sharing
 - ❑ P_1 , P_2 , and P_3 share mailbox A
 - ❑ P_1 , sends; P_2 and P_3 receive
 - ❑ Who gets the message?
- ❑ Solutions
 - ❑ Allow a link to be associated with at most two processes
 - ❑ Allow only one process at a time to execute a receive operation
 - ❑ Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was. (Algorithm like RR)

Blocking and Non-blocking schemes

- ❑ Message passing may be either blocking or non-blocking
- ❑ **Blocking** is considered **synchronous**
 - ❑ **Blocking send** -- the sender is blocked until the message is received
 - ❑ **Blocking receive** -- the receiver is blocked until a message is available
- ❑ **Non-blocking** is considered **asynchronous**
 - ❑ **Non-blocking send** -- the sender sends the message and continue
 - ❑ **Non-blocking receive** -- the receiver receives:
 - ❑ A valid message, or
 - ❑ Null message
- ❑ Different combinations possible
 - ❑ If both send and receive are blocking, we have a **rendezvous**

Buffering

- ❑ Whether communication is direct or indirect, messages exchanged by communicating processes reside in a temporary queue.
- ❑ Such queues can be implemented in three ways:
 - ❑ 1. Zero capacity – no messages are queued on a link. Sender must wait for receiver (rendezvous)
 - ❑ 2. Bounded capacity – finite length of n messages Sender must wait if link full
 - ❑ 3. Unbounded capacity – infinite length Sender never waits

Example of IPC Systems

- ❑ There are four different IPC systems.
 - ❑ POSIX API for shared memory
 - ❑ Mach operating system, which uses message passing
 - ❑ Windows IPC, which uses shared memory as a mechanism for providing certain types of message passing.
 - ❑ Pipes, one of the earliest IPC mechanisms on UNIX systems.

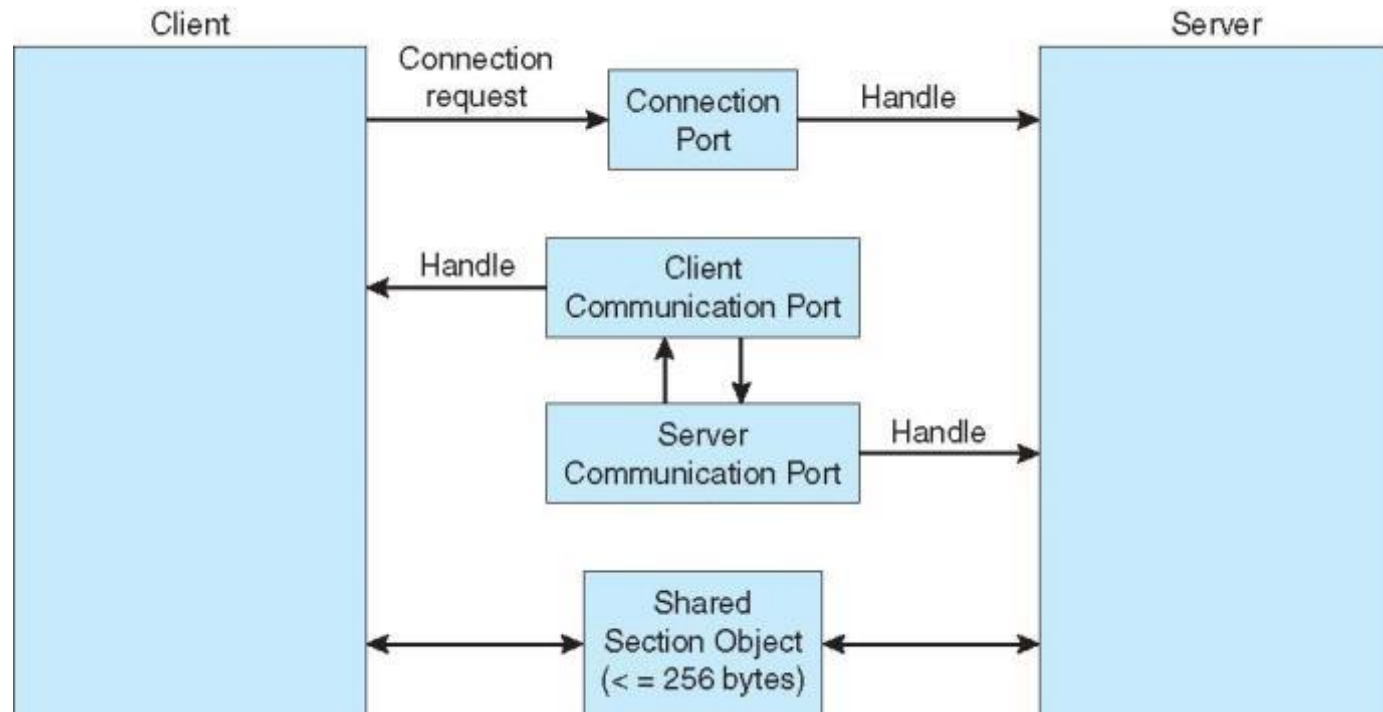
Mach

- ❑ Mach communication is message based
 - ❑ Even system calls are messages
 - ❑ Each task gets two mailboxes at creation- Kernel and Notify
 - ❑ Only three system calls needed for message transfer
 - ❑ `msg_send()`, `msg_receive()`, `msg_rpc()`
 - ❑ Mailboxes needed for communication, created via
 - ❑ `port_allocate()`
 - ❑ Send and receive are flexible, for example four options if mailbox full:
 - ❑ Wait indefinitely
 - ❑ Wait at most n milliseconds
 - ❑ Return immediately
 - ❑ Temporarily cache a message

Windows

- ❑ Message-passing centric via **advanced local procedure call (LPC)** facility
 - ❑ Only works between processes on the same system
 - ❑ Uses ports (like mailboxes) to establish and maintain communication channels
 - ❑ Communication works as follows:
 - ❑ The client opens a handle (an abstract reference to a resource) to the subsystem's **connection port** object.
 - ❑ The client sends a connection request.
 - ❑ The server creates a private **communication port** and returns the handle to the client.
 - ❑ The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.

Local Procedure Calls in Windows



Communications in Client-Server Systems

- ☐ Sockets
- ☐ Remote Procedure Calls
- ☐ Remote Method Invocation (Java)



Books

- ❑ Operating Systems Concept
 - ❑ Written by Galvin and Silberschatz
 - ❑ Edition: 9th



References

- ❑ Operating Systems Concept
 - ❑ Written by Galvin and Silberschatz
 - ❑ Edition: 9th